

Ewan John Dennis

AI/ML Engineer | Software Developer

+91 7994687722 | ewandennis2006@gmail.com | linkedin.com/in/ewanjohndennis |
github.com/ewanjohndennis | Kochi, Kerala

SUMMARY

Computer Science Engineering student with hands-on experience building agentic AI systems, ML pipelines, and data-intensive backend systems. Shipped a production multi-agent market intelligence platform during my Infosys Springboard internship and currently building a content opportunity scoring model at FlyRank AI against a 79M-row data warehouse. Work well independently, iterate quickly, and care about systems that hold up outside controlled environments.

EXPERIENCE

FlyRank AI

Jul 2026 – Present

Machine Learning Intern

Remote

- Building a content refresh prioritization model (Applied Search Intelligence track) that ranks which web pages a reviewer should look at first, across a 519K-item content warehouse with ~79M daily performance rows.
- Benchmarked a learned random forest model against FlyRank's existing rule-based system: Precision@50 improved from 0.240 to 0.740 on a 30,000-page starter dataset — 37 correct picks in the top 50 vs 12 for the baseline.
- Identified search visibility consistency (`days_with_impressions`) as the top predictive feature, ahead of content age or word count. Applied client-holdout validation to ensure the model generalizes to unseen clients, not just unseen pages.
- Designing a future-window label: using 90 days of prior signals to predict content decline over the next 30 days, replacing the current proxy label. Stack: Python, scikit-learn, DuckDB, pandas, Hugging Face datasets.

Infosys Springboard (Virtual Internship 6.0)

Feb 2026 – Apr 2026

AI Engineer Intern

Remote

- Designed and deployed a multi-agent AI architecture with three parallel agents running concurrently via `ThreadPoolExecutor`, reducing data-collection latency by up to 3x compared to sequential execution.
- Built a synthesis pipeline integrating a FAISS vector database for RAG, a Groq LLM for report generation, and a PyTorch LSTM for time-series forecasting, running end-to-end with no manual steps.
- Iterated on prompt engineering pipelines across Azure OpenAI, HuggingFace, OpenRouter, and Groq — evaluating output quality and reliability under real-world rate limits and tracking provider performance across runs.
- Shipped a role-based Streamlit UI with MongoDB Atlas persistence, Gmail SMTP delivery, and PDF reporting via ReportLab. Generated 20+ intelligence reports in production.

PROJECTS

RTIIS | Real-Time Industry Insight & Strategic Intelligence System

2026

- Flagship project from Infosys Springboard Virtual Internship 6.0. Multi-agent platform for automated market intelligence: three parallel agents collect news, competitor data, and financial signals concurrently via `ThreadPoolExecutor`, cutting data-collection latency by up to 3x.
- Synthesis agent uses FAISS for RAG, Groq for report generation, and a PyTorch LSTM for stock trend forecasting. Deployed on HuggingFace Spaces with MongoDB Atlas persistence, Gmail SMTP delivery, and PDF export via ReportLab. Supports up to 4 competitor companies per run.
- **GitHub:** github.com/Ewanjohndennis/RealTimeMarketIntelli

SentinelAI | AI-Enhanced DevOps Assistant

2026

- Built the AI reasoning engine and safety validation layer for a 4-person DevOps assistant. The safety layer uses a 0.6 confidence threshold and 120-second cooldown before executing infrastructure actions, preventing unsafe autonomous operations.
- Powered by Groq's LLaMA-3.3-70b-versatile via FastAPI, with Prometheus for system monitoring and Docker for containerisation. Submitted to Project Jam hackathon (June 2026).
- **GitHub:** github.com/AlvinGeorge-AG/AI-Enhanced-DevOps-Assistant

- Published to PyPI. Trains scikit-learn and XGBoost models directly from a CSV file with a single command, covering preprocessing, training, evaluation, and export. XGBoost is a lazy conditional dependency to keep the base install lightweight.
- Ships with an MCP server component, exposing training and inference as tools callable by any MCP-compatible agent.
- **GitHub:** github.com/Ewanjohndennis/traincli

- Real-time platform for students to find and join active study sessions on campus. Validated real-time sync under concurrent sessions with sub-100ms Firestore latency; auto-cleans expired sessions. **Awarded 2nd Prize at Techsprint (GDGC MEC)**, from ~40 teams, selected into a 10-team final.
- **GitHub:** github.com/Ewanjohndennis/huddle **Live:** huddlechats.vercel.app

SKILLS

AI / ML: Transformer Architecture, Agentic Workflows, Tool Calling, Embedding Models, Retrieval-Augmented Generation (RAG), Multi-Agent Systems, LLM Integration, Context Engineering, FAISS Vector Database, Vector Embeddings, PyTorch, LSTM Forecasting, NLP, Fine-Tuning, Model Evaluation, MLOps, Data Preprocessing, scikit-learn, XGBoost, DuckDB

Providers & SDKs: Groq, HuggingFace Inference API, OpenRouter, Azure OpenAI, Anthropic SDK

Backend: Python, Node.js, Express.js, Spring Boot, REST API, PostgreSQL, MongoDB, Docker, FastAPI

Frontend: JavaScript (ES6+), TypeScript, React.js, Next.js, Tailwind CSS, Firebase

DevOps / Cloud: Git, CI/CD, Google Cloud Run, Vercel, Streamlit Cloud, HuggingFace Spaces

Methodologies: Agile (Scrum), Data Structures & Algorithms

CERTIFICATIONS

AI Fluency: Framework & Foundations — Anthropic	Feb 2026
Introduction to Generative AI — Google Cloud	Oct 2025
Artificial Intelligence Fundamentals — IBM	Jun 2025
Postman API Fundamentals Student Expert — Postman	May 2025

EDUCATION

Model Engineering College	2024 – Present
<i>B.Tech in Computer Science and Engineering — CGPA: 8.73</i>	Kochi, Kerala
Vidyodaya School, Thevakkal	2022 – 2024
<i>Senior Secondary (Class 12) — Percentage: 93%</i>	CBSE